

# Offensive Language and Sentiment Analysis for User Comments

## Table of Contents

1. Background .....	1
2. Requirements .....	1
3. Method .....	2
4. Implementation .....	2
5. Milestones .....	3
6. Gathering Results .....	3

## 1. Background

This project aims to implement a system for analyzing user comments, focusing on sentiment analysis, offensive language detection, and categorization. The objective is to provide a nuanced classification of comments, identifying offensive content while refining overall sentiment scoring to allow for better moderation and filtering of inappropriate or harmful messages.

The initial implementation used a basic sentiment analysis model and iteratively introduced various offensive language detection techniques. These enhancements aimed to provide contextually aware categorizations and accommodate varying degrees of offensiveness.

## 2. Requirements

- **Must Have:**

- Load and preprocess user comments in a standardized format.
- Perform sentiment analysis on user comments using a pre-trained model.
- Detect and classify comments containing offensive language.
- Implement a weighted sentiment adjustment for comments flagged as offensive.

- **Should Have:**

- Handle different comment lengths, ensuring compatibility with model input size constraints.
- Provide granular categorization for comments based on offensive word density and patterns.
- Maintain a modular design for easy modification and integration of additional features.

- **Could Have:**

- Support multi-language sentiment and offensive language analysis.

- Use context-aware models to further enhance classification accuracy.
- Implement a sarcasm detection mechanism for refining sentiment scores.

## 3. Method

The project is divided into two primary phases:

1. **Sentiment Analysis Integration:** The `distilbert-base-uncased-finetuned-sst-2-english` model was used for initial sentiment classification. Key steps include:
  - Environment setup and import of necessary libraries (`transformers`, `torch`, `pandas`).
  - Data loading and preprocessing with error handling for malformed entries.
  - Batch processing for efficient handling of large datasets.
2. **Offensive Language Detection and Weight Adjustment:** Offensive language detection is performed in multiple stages, focusing on:
  - Flagging highly offensive words using a predefined list.
  - Implementing word density-based weighting adjustments.
  - Fine-tuning offensive categories with context-based severity modifiers.

## 4. Implementation

1. **Data Preparation and Setup (Cells 1-2):**
  - **Environment Setup and Imports:** Initialize libraries and define helper functions (`refine_labels`).
  - **Dataset Loading:** Handle file errors and verify the presence of required columns.
  - Extract comments from the dataset and convert them to a standardized string format for analysis.
2. **Sentiment Analysis (Cell 3):**
  - Use the pre-trained `distilbert-base-uncased-finetuned-sst-2-english` model.
  - Process comments in batches to optimize memory usage.
  - Store the results in a structured format (`df_sentiments` DataFrame) for further analysis.
3. **Sentiment Distribution Analysis (Cell 4):**
  - Analyze the distribution of positive, negative, and neutral sentiments.
  - Calculate average sentiment scores for each category.
4. **Extreme Sentiment Extraction (Cell 5):**
  - Identify and display the top 5 positive and negative comments based on sentiment scores.
5. **Offensive Language Detection (Cell 6):**
  - Implement a word frequency analysis using regex-based tokenization.
  - Identify the top 50 most common words for reference.

#### 6. Flagging Highly Offensive Comments (Cell 7):

- Create a custom list for words of interest (e.g., racial slurs).
- Add an **Offensive\_Flag** column to flag comments with highly offensive terms.
- Update the **Sentiment Label** for flagged comments to **Highly Offensive**.

#### 7. Sentiment Label Adjustment (Cell 8):

- Modify sentiment labels for flagged comments.
- Adjust sentiment scores to indicate the intensity of offensive language.

#### 8. Summary Report (Cell 9):

- Provide a text-based summary of sentiment distribution and offensive language identification.
- Highlight key insights and findings from the sample analysis.

## 5. Milestones

#### 1. Milestone 1: Environment Setup and Data Loading:

- Complete the environment configuration.
- Verify dataset integrity and perform initial extraction.

#### 2. Milestone 2: Sentiment Analysis Integration:

- Implement sentiment analysis using **transformers** pipelines.
- Validate the model output for a sample of comments.

#### 3. Milestone 3: Offensive Language Detection:

- Develop and test offensive language detection using predefined word lists.
- Implement and validate the offensive language flagging mechanism.

#### 4. Milestone 4: Sentiment Label Refinement and Weight Adjustments:

- Refine labels based on detected offensive language.
- Finalize scoring mechanisms and evaluate output consistency.

## 6. Gathering Results

The evaluation phase includes running the detection on a comprehensive set of sample comments and validating the output accuracy. Metrics for evaluation include:

- **Accuracy of Sentiment Classification:** Measure the consistency between predicted and actual sentiment.
- **Precision of Offensive Language Detection:** Evaluate the percentage of correctly flagged comments.
- **Weighted Sentiment Score Impact:** Analyze how offensive flags influence overall sentiment scoring.

The refined model demonstrates a balanced approach to comment classification, accounting for both sentiment intensity and the presence of harmful language. Future enhancements could include context-aware analysis and support for a broader range of languages.